

Kiefer-Wolfowitz Thm & G-optimal / D-optimal Designs.

Source: Chapter 21, Bandit Algorithms, L & S. (textbook)

Previously: Choose $\{A_s, s=1, 2, \dots, t\}$ sequentially in \mathbb{R}^d according to LinUCB rule. Develop $\{C_s, s=1, 2, \dots, t\}$ confidence sets for $\{\hat{\theta}_s, s=1, \dots, t\}$ associated with the regularized least-square estimator.

A minmax problem — We are given a compact set $A \subseteq \mathbb{R}^d$ (set could be finite/infinite). Let us define π to be a probability distribution over any finite subset of A , denoted as $\text{supp}(\pi)$. Let

$$g(\pi) = \max_{a \in A} \|a\|_V(\pi)^{-1},$$

$$\text{where } V(\pi) = \sum_{a \in A} \pi(a) a a^\top.$$

Then, objective is to find a π^* s.t.

$$\pi^* \in \arg \min_{\pi} g(\pi).$$

(π^* is called as the G-optimal Design ...)

$\boxed{\begin{array}{l} \text{A maximization problem} \\ \pi \text{ a dist. over } \Delta \subseteq \mathbb{R}^d \end{array}}$ — Same setting as above, with

Let $f(\pi) = \log \det(\nabla(\pi))$.

Then, objective is to find $\hat{\pi}$ s.t.

$$\boxed{\hat{\pi} \in \underset{\pi}{\operatorname{argmax}} f(\pi)} \quad (\hat{\pi} \text{ is called as the D-optimal Design})$$

Least Squares Summary: Recall that:

$$y_i = \langle A_i, \theta^* \rangle + \eta_i, \quad (i=1, 2, \dots, n)$$

with $\{A_1, \dots, A_n\} \in \mathbb{R}^d$, θ^* in \mathbb{R}^d , the following holds from Linear Least Squares Regression:

$$A = \begin{pmatrix} & \xleftarrow{d} \\ A_1^T & \\ A_2^T & \\ \vdots & \\ A_n^T & \end{pmatrix} \quad \begin{matrix} \uparrow \\ n \\ \downarrow \end{matrix} \quad \begin{matrix} \gamma \\ \vdots \\ \gamma_n \end{matrix}^T \quad \begin{matrix} \xrightarrow{n} \\ \eta \\ \eta_n \end{matrix}^T$$

$\gamma = A \theta^* + \eta$

\mathbb{R}^n

design matrix

$$\{A_i\}_{i=1}^n \text{ spans } \mathbb{R}^d.$$

Then, $\hat{\theta} = V^{-1} \sum_{i=1}^n A_i X_i$, unbiased estimate, where

$$V = \sum_{i=1}^n A_i A_i^\top = (A^\top A)$$

Further V^{-1} is the covariance matrix, i.e.,

$$E \left[(\hat{\theta} - \theta^*) (\hat{\theta} - \theta^*)^\top \right] = V^{-1}$$

Further along any vector $x \in \mathbb{R}^d$, the variance is

$$E \left[\langle \hat{\theta} - \theta^*, x \rangle^2 \right] = x^\top V^{-1} x = \|x\|_V^2$$

G-optimal Design

D-optimal Design

$$\pi^* \in \arg \min_{\pi} \max_{a \in A} \|a\|_{V(\pi)^{-1}}^2$$

$$\hat{\pi} \in \arg \max_{\pi} \log \det(V(\pi))$$

p.d. symmetric
matrix

Intuition: Equalize the variance along all "core" vectors in A .

Intuition: $\det(V(\pi))$ is the product of eigenvalues. This scales inversely with the volume of the confidence ellipsoid around $(\hat{\theta} - \theta^*)$.

Ref: Optimal Design, Notes by Steven Buyske & Richard Trout,
Rutgers University (Class notes for Stat. 591)

Geometry : Let π^* solve the problem above, and the corresponding $V(\pi^*) = \sum_{a \in A} \pi^*(a) a a^\top$.

Then, $(g(\pi^*), V(\pi^*))$ defines the minimum volume centered ellipsoid containing A . More concretely :

$$\text{Let } \Sigma = \left\{ x \in \mathbb{R}^d : \|x\|_{V(\pi^*)^{-1}}^2 \leq g(\pi^*) \right\}.$$

$= d$ (we will show later).

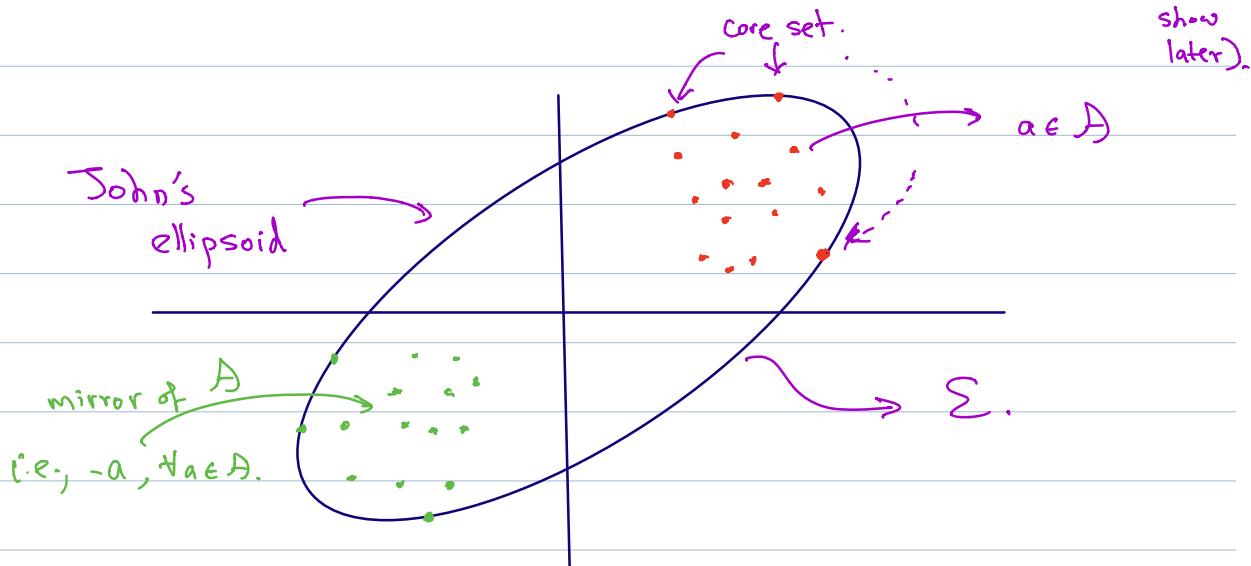


Figure 21.1 in Bandit Algorithms (textbook)

Then Σ is the minimum volume centered (at zero) ellipsoid containing the set A .

Why is this useful? π^* a G-optimal Design for A , and has finite support $\text{supp}(A)$.

Now, consider the following design problem: We are allowed to pre-decide an integer $n \geq 1$, and a sequence of vectors $\{a_1, a_2, \dots, a_n\}$, $a_i \in A$. There is some unknown $\theta^* \in \mathbb{R}^d$, and we get the usual feedback, i.e., we observe:

$$x_i = \langle a_i, \theta^* \rangle + \eta_i, \quad \eta_i \sim \text{1-subG}.$$

We construct the least-square estimator $\hat{\theta}$, where

$$\hat{\theta} = V^{-1} \sum_{i=1}^n a_i x_i, \quad \text{with } V = \sum_{i=1}^n a_i a_i^\top$$

Objective: Given (ε, δ) , $\varepsilon > 0$, $\delta \in (0, 1)$ and any $a \in A$, we want

$$P(\langle \hat{\theta} - \theta^*, a \rangle \geq \varepsilon) \leq \delta$$

function of ...

Given (ε, δ) , determine the smallest n , and associated $\{a_1, a_2, \dots, a_n\}$ so that above bound holds.

The G-optimal Design provides an upper bound for n and a construction for $\{a_1, \dots, a_n\}$. Textbook uses the term "accurate approximation".

Construction: π^* a G-optimal Design for A , and
 $\text{Supp}(\pi^*)$ finite (notation is slightly bad, as it should be
 $\text{Supp}_A(\pi^*)$, as the support is a function of A).

Let $n_a = \left\lceil \frac{2\pi^*(a) g(\pi^*)}{\varepsilon^2} \ln\left(\frac{1}{\delta}\right) \right\rceil$,

$\forall a \in \text{Supp}(A)$

(Aside: Theorem below will show that $|\text{Supp}(A)| \leq \frac{d(d+1)}{2}$)

Let $n = \sum_{a \in \text{Supp}(\pi^*)} n_a$. We deterministically construct

$\{a_1, a_2, \dots, a_n\}$, by replicating $a \in \text{Supp}(\pi^*)$ n_a times, $\forall a \in \text{Supp}(\pi^*)$.

Thus, $V = \sum_{i=1}^n a_i a_i^\top = \sum_{a \in \text{Supp}(\pi^*)} n_a a a^\top$

$$\geq \sum_{a \in \text{Supp}(\pi^*)} 2 \cdot \underbrace{\pi^*(a) g(\pi^*)}_{\varepsilon^2} \ln\left(\frac{1}{\delta}\right) \cdot \underbrace{a a^\top}_{\text{rank } 1}$$

$$= \frac{2g(\pi^*)}{\varepsilon^2} \ln\left(\frac{1}{\delta}\right) V(\pi^*) \rightarrow \text{XX}$$

Recall that from Set 15, for any $a \in \mathbb{R}^d$, and $\delta \in (0, 1)$,

$$P\left(\langle \hat{\theta} - \theta^*, a \rangle \geq \sqrt{2 \|a\|_{V^{-1}}^2 \ln(1/\delta)}\right) \leq \delta.$$

From $(**)$, we have for any $a \in A$,

$$\|a\|_{V^{-1}}^2 = a^\top V^{-1} a \leq (a^\top V(\pi^*)^{-1} a) \frac{\varepsilon^2}{2 g(\pi^*) \ln(1/\delta)}.$$

$$\leq \frac{\varepsilon^2}{2 \ln(1/\delta)}.$$

$$\Rightarrow \sqrt{2 \ln(1/\delta) \|a\|_{V^{-1}}^2} \geq \varepsilon$$

$$\therefore P\left(\langle \hat{\theta} - \theta^*, a \rangle \geq \varepsilon\right) \leq \delta$$

$$\text{Further: } n = \sum_{a \in \text{Supp}(\pi^*)} n_a = \sum_{a \in \text{Supp}(\pi^*)} \left\lceil \frac{2 \pi^*(a) g(\pi^*) \ln(1/\delta)}{\varepsilon^2} \right\rceil$$

$$\leq \underbrace{|\text{Supp}(\pi^*)|}_{\text{core set}} + \frac{2g(\pi^*)}{\varepsilon^2} \ln(1/\delta).$$

$$\text{core set} = \{a \in A : a \in \text{Supp}(\pi^*)\}$$

Kiefer-Wolfowitz Thm (21.1 in text): Suppose $A \subseteq \mathbb{R}^d$, is a compact set with $\text{Span}(A) = \mathbb{R}^d$. Then,

$$(i) \pi^* \in \operatorname{argmin}_{\pi} g(\pi) \iff (ii) g(\pi^*) = d.$$



$$(ii) \pi^* \in \operatorname{argmax}_{\pi} \log \det(V(\pi))$$

Further $\exists \pi^*$ as above with $\text{Supp}(\pi^*) \leq \frac{d(d+1)}{2}$

Remark: ① $\pi^* \in \operatorname{argmin}_{\pi} g(\pi) \rightarrow G\text{-Optimal Design}$

② $\pi^* \in \operatorname{argmax}_{\pi} \log \det(V(\pi)) \rightarrow D\text{-Optimal Design.}$

Above Thm shows these are equivalent (essentially by establishing that the problems are duals of each other).

Proof: (Spcl case, where $|A|$ is finite). Thus, $\pi \in \mathbb{R}^{|A|}$, a finite dim. vector over A . Recall:

$$g(\pi) = \max_{a \in A} \|a\|_{V(\pi)}^2 \quad f(\pi) = \log \det V(\pi).$$

Facts: f is concave, and $\nabla f(\pi) = \begin{pmatrix} \vdots \\ \|a\|_{V(\pi)}^2 \\ \vdots \end{pmatrix}_{|A|}$

Also, for any π , we have

$$\begin{aligned} \sum_{a \in A} \pi(a) \|a\|_{V(\pi)}^2 &= \sum_{a \in A} \pi(a) \cdot a^\top V(\pi)^{-1} a \\ &= \text{trace} \left(\underbrace{\sum_a \pi(a) a a^\top V(\pi)^{-1}}_{V(\pi)} \right) \\ &= \text{trace}(I) = d. \end{aligned}$$

$$\therefore g(\pi) = \max_{a \in A} \|a\|_{V(\pi)}^2 \geq \underbrace{\sum_{a \in A} \pi(a) \|a\|_{V(\pi)}^2}_{\text{max} \geq \text{avg}} = d.$$

$\hookrightarrow \textcircled{I}$

$$\textcircled{2}. \quad \pi^* \in \arg \max_{\pi} \underbrace{\log \det V(\pi)}_{f(\pi)} \implies \pi^* \in \arg \min_{\pi} \max_{a \in A} \underbrace{\|a\|_{V(\pi)}^2}_{g(\pi)}.$$

f concave. From first-order optimality conditions, for any π , we have

$$0 \geq \langle \nabla f(\pi^*), \pi - \pi^* \rangle$$

$$= \sum_{a \in A} \pi(a) \|a\|_{V(\pi^*)^{-1}}^2 - \underbrace{\sum_{a \in A} \pi^*(a) \|a\|_{V(\pi^*)^{-1}}^2}_{= d \text{ (from (i))}}$$

Now choose any $\tilde{a} \in A$, and let

$$\pi(\tilde{a}) = 1, \quad \pi(a) = 0 \quad \forall a \neq \tilde{a}.$$

$$\text{Then, } 0 \geq \|\tilde{a}\|_{V(\pi^*)^{-1}}^2 - d$$

$$\Rightarrow \|\tilde{a}\|_{V(\pi^*)^{-1}}^2 \leq d \quad \text{for any } \tilde{a} \in A$$

$$\Rightarrow \max_{a \in A} \|a\|_{V(\pi^*)^{-1}}^2 \leq d$$

$$\Rightarrow g(\pi^*) \leq d$$

$$\text{Also, } g(\pi) \geq d \quad \forall \pi \quad (\text{from (i)}).$$

$\Rightarrow \pi^*$ is a minimizer for $g(\pi)$, and
 $g(\pi^*) = d$.

(b) $g(\pi^*) = d \Rightarrow \pi^* \in \operatorname{argmax}_\pi f(\pi)$.

We have $g(\pi^*) = \max_{a \in A} \|a\|_{V(\pi^*)^{-1}}^2 = d$

Now observe that:

$$\langle \nabla f(\pi^*), \pi - \pi^* \rangle$$

$$= \sum_{a \in A} \pi(a) \cdot \|a\|_{V(\pi^*)^{-1}}^2 - \underbrace{\sum_{a \in A} \pi^*(a) \|a\|_{V(\pi^*)^{-1}}^2}_{\leq d} = d.$$

$\forall a \in A$

$$\therefore \max_{a \in A} \|a\|_{V(\pi^*)^{-1}}^2 = d.$$

$$\leq 0$$

$$f \text{ concave}, \quad \langle \nabla f(\pi^*), \pi - \pi^* \rangle \leq 0$$

$\Rightarrow \pi^*$ is the maximizer of f .

(c) $\pi^* \in \arg\min_{\pi} \max_{a \in A} \|a\|_{V(\pi)^{-1}}^2 \Rightarrow g(\pi^*) = d$.

Recall $g(\pi) \geq d$ $\forall \pi$ (from ②).

We now demonstrate a $\hat{\pi}$ s.t. $g(\hat{\pi}) = d$. Thus, it immediately follows that $g(\pi^*) = d$.

Let $\hat{\pi}$ be the maximizer of $f(\cdot)$, i.e.,

$$\hat{\pi} \in \operatorname{argmax}_{\pi} \log \det V(\pi).$$

From (a), we note $g(\hat{\pi}) = d$. Thus, we are done.

For showing that $\exists \pi^* \text{ s.t. } \operatorname{Supp}(\pi^*) \leq \frac{d(d+1)}{2}$,

read text, Pf & Thm 21.1

